

Índice

- 1 MultinomialNB
- 2 LinearSVC
- 3 Logistic Regression

```
In [1]: from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import LinearSVC
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
from sklearn.multiclass import OneVsRestClassifier

import pickle
X_train_smt = pickle.load(open("saved_feats/ShuffleSplit_X_train_smt", "rb"))
X_test = pickle.load(open("saved_feats/ShuffleSplit_X_test", "rb"))
y_train_smt = pickle.load(open("saved_feats/ShuffleSplit_y_train_smt", "rb"))
y_test = pickle.load(open("saved_feats/ShuffleSplit_y_test", "rb"))

In [2]: param_grid_mnb = {'alpha': [1, 0.1, 0.01, 0.001]}
grid_mnb = GridSearchCV(MultinomialNB(fit_prior=True,
                                     class_prior=[0.2, 0.5, 0.93]), param_grid_mnb)
print("Mejor modelo predicho: " + str(grid_mnb.best_estimator_))

Mejor modelo predicho: MultinomialNB(alpha=0.001, class_prior=[0.2, 0.5, 0.93])
```

LinearSVC

```
In [3]: param_grid_svc = {'C': [0.1, 1, 10, 100]}
grid_svc = GridSearchCV(LinearSVC(penalty='l1', loss='squared_hinge', dual=False, class_weight='Negative'),
                        param_grid_svc, scoring='accuracy', refit=True, verbose=2)
grid_svc.fit(X_train_smt, y_train_smt)
print("Mejor modelo predicho: " + str(grid_svc.best_estimator_))

Fitting 5 folds for each of 4 candidates, totalling 20 fits
[CV] END .....C=0.1; total time= 2.6s
[CV] END .....C=0.1; total time= 2.6s
[CV] END .....C=0.1; total time= 2.6s
[CV] END .....C=0.1; total time= 2.7s
[CV] END .....C=0.1; total time= 2.8s
[CV] END .....C=1; total time= 15.7s
[CV] END .....C=1; total time= 16.0s
[CV] END .....C=1; total time= 16.6s
[CV] END .....C=1; total time= 17.5s
[CV] END .....C=1; total time= 18.9s
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1225: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
[CV] END .....C=10; total time= 25.4s
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1225: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(
[CV] END .....C=10; total time= 25.1s
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1225: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(
[CV] END .....C=10; total time= 24.8s
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1225: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(
[CV] END .....C=10; total time= 25.1s
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1225: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(
[CV] END .....C=10; total time= 25.5s
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1225: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(
[CV] END .....C=100; total time= 21.6s
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1225: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(
[CV] END .....C=100; total time= 21.5s
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1225: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(
[CV] END .....C=100; total time= 21.5s
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1225: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(
[CV] END .....C=100; total time= 21.6s
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1225: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(
[CV] END .....C=100; total time= 22.3s
Mejor modelo predicho: LinearSVC(C=10,
                                class_weight={'Negative': 0.2, 'Neutral': 0.5, 'Positive': 0.93},
                                dual=False, penalty='l1')
/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py:1225: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(

In [4]: param_grid_svc = {'loss': ['hinge', 'squared_hinge'], 'C': [0.1, 1, 10, 100]}
grid_svc = GridSearchCV(LinearSVC(penalty='l2', dual=False, class_weight={'Negative': 0.2, 'Neutral': 0.5, 'Positive': 0.93},
                                scoring='accuracy', refit=True, verbose=2)
                        param_grid_svc, scoring='accuracy', refit=True, verbose=2)
grid_svc.fit(X_train_smt, y_train_smt)
print("Mejor modelo predicho: " + str(grid_svc.best_estimator_))

Fitting 5 folds for each of 8 candidates, totalling 40 fits
[CV] END .....C=0.1, loss=hinge; total time= 0.1s
[CV] END .....C=0.1, loss=hinge; total time= 0.1s
[CV] END .....C=0.1, loss=hinge; total time= 0.1s
[CV] END .....C=0.1, loss=hinge; total time= 0.1s
[CV] END .....C=0.1, loss=squared_hinge; total time= 0.7s
[CV] END .....C=0.1, loss=squared_hinge; total time= 0.7s
[CV] END .....C=0.1, loss=squared_hinge; total time= 0.7s
[CV] END .....C=0.1, loss=hinge; total time= 0.1s
[CV] END .....C=1, loss=hinge; total time= 0.1s
[CV] END .....C=1, loss=hinge; total time= 0.1s
[CV] END .....C=1, loss=hinge; total time= 0.1s
[CV] END .....C=1, loss=squared_hinge; total time= 1.3s
[CV] END .....C=1, loss=squared_hinge; total time= 1.1s
[CV] END .....C=1, loss=squared_hinge; total time= 1.2s
[CV] END .....C=10, loss=hinge; total time= 0.1s
[CV] END .....C=10, loss=hinge; total time= 0.1s
[CV] END .....C=10, loss=squared_hinge; total time= 2.6s
[CV] END .....C=10, loss=squared_hinge; total time= 2.6s
[CV] END .....C=10, loss=squared_hinge; total time= 3.3s
[CV] END .....C=100, loss=hinge; total time= 10.7s
[CV] END .....C=100, loss=hinge; total time= 0.1s
[CV] END .....C=100, loss=hinge; total time= 0.1s
[CV] END .....C=100, loss=squared_hinge; total time= 7.2s
[CV] END .....C=100, loss=squared_hinge; total time= 10.7s
[CV] END .....C=100, loss=squared_hinge; total time= 7.3s
[CV] END .....C=100, loss=squared_hinge; total time= 7.9s
[CV] END .....C=100, loss=squared_hinge; total time= 4.8s
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:
20 fits failed out of a total of 40.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:
-----
20 fits failed with the following error:
Traceback (most recent call last):
  File "/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/usr/local/lib/python3.10/dist-packages/sklearn/svm/_classes.py", line 257, in fit
    self.coef_, self.intercept_, n_iter_ = fit_liblinear(
  File "/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py", line 1204, in _fit_liblinear
    solver_type = get_liblinear_solver_type(multi_class, penalty, loss, dual)
  File "/usr/local/lib/python3.10/dist-packages/sklearn/svm/_base.py", line 1043, in _get_liblinear_solver_type
    raise ValueError(
ValueError: Unsupported set of arguments: The combination of penalty='l2' and loss='hinge' are not supported when dual=False, Parameters: penalty='l2', loss='hinge', dual=False

warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:953: UserWarning: One or more of the test scores are non-finite: [ nan 0.92701578 nan 0.95855072 nan 0.96852022 nan 0.96947898]
  warnings.warn(
Mejor modelo predicho: LinearSVC(C=100,
                                class_weight={'Negative': 0.2, 'Neutral': 0.5, 'Positive': 0.93},
                                dual=False)
```

Logistic Regression

```
In [10]: param_grid_logit = {'penalty': ['l1', 'l2', 'elasticnet', 'none'], 'C': [0.1, 1, 10, 100],
                        'solver': ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga']}
grid_logit = GridSearchCV(LogisticRegression(class_weight={'Negative': 0.2, 'Neutral': 0.5, 'Positive': 0.93},
                                             param_grid_logit, verbose=2)
                        param_grid_logit, verbose=2)
grid_logit.fit(X_train_smt, y_train_smt)
```


[illegible]

[illegible]

