

Hyperparameters Logistic Regression

Índice

- 1 **Parámetros por defecto**
- 2 **Penalty=none**
- 3 **Solver=newton-cg**
- 4 **Penalty=none, solver=newton-cg**
- 5 **Penalty=l1, solver=liblinear, multi_class=ovr**
- 6 **Dual=True, solver=liblinear, multi_class=ovr**
- 7 **Solver=sag**
- 8 **Penalty=none, solver=sag**
- 9 **Penalty=elasticnet, solver=saga, l1_ratio=0.5**
- 10 **Penalty=l1, solver=saga**
- 11 **Solver=saga**
- 12 **Penalty=none**

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import scipy.stats as st

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report, roc_curve, roc_auc_score
from sklearn.multiclass import OneVsRestClassifier
from sklearn.model_selection import train_test_split

import pickle

X_train_smt = pickle.load(open("saved_feats/ShuffleSplit_X_train_smt", "rb"))
X_test = pickle.load(open("saved_feats/ShuffleSplit_X_test", "rb"))
y_train_smt = pickle.load(open("saved_feats/ShuffleSplit_y_train_smt", "rb"))
y_test = pickle.load(open("saved_feats/ShuffleSplit_y_test", "rb"))

# FORMATO DE PLOTS
plt.style.use('bmh')
```

Parámetros por defecto

```
In [2]: clf_logit = LogisticRegression(penalty='l2', dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1,
random_state=None, solver='lbfgs', max_iter=1000, multi_class='auto', verbose=0)
pred_logit = clf_logit.predict(X_test)
pred_prob_logit = clf_logit.predict_proba(X_test)
```

```
In [3]: print('Predicted classes:', clf_logit.classes_)
print('Average accuracy:', np.mean(pred_logit == y_test)*100)
print('Train accuracy:', (clf_logit.score(X_train_smt, y_train_smt))*100)
print('Test accuracy:', (clf_logit.score(X_test, y_test))*100)

print('\n CONFUSION MATRIX')
print(confusion_matrix(y_test, pred_logit))
print('\n CLASSIFICATION REPORT')
print(classification_report(y_test, pred_logit))

Predicted classes: ['Negative' 'Neutral' 'Positive']
Average accuracy: 87.4701451405475
Train accuracy: 95.7668045207964
Test accuracy: 87.4701451405475

CONFUSION MATRIX
[[ 67  24  28]
 [ 48 104  98]
 [ 96 388 458]]

CLASSIFICATION REPORT
              precision    recall  f1-score   support

Negative      0.32         0.56         0.41        119
Neutral       0.20         0.42         0.27        250
Positive      0.97         0.90         0.94       5074

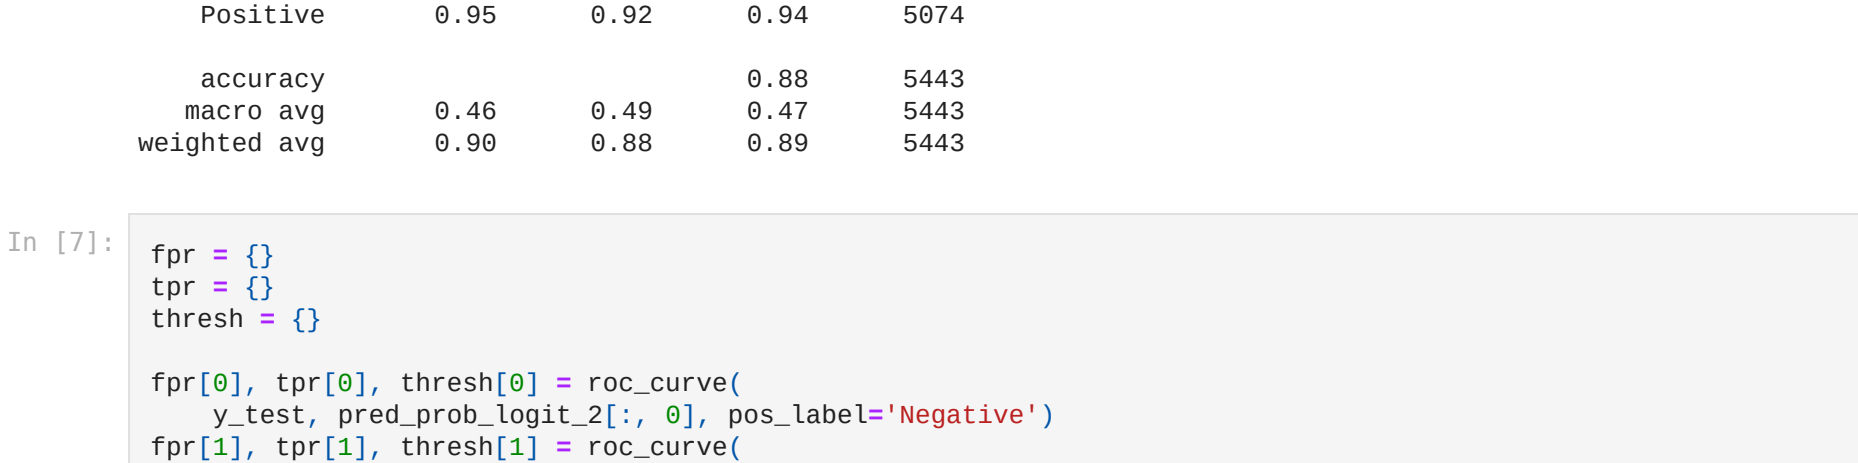
accuracy      0.50         0.63         0.87       5443
macro avg     0.50         0.63         0.54       5443
weighted avg   0.92         0.87         0.90       5443
```

```
In [4]: fpr = {}
tpr = {}
thresh = {}

fpr[0], tpr[0], thresh[0] = roc_curve(y_test, pred_prob_logit[:, 0], pos_label='Negative')
fpr[1], tpr[1], thresh[1] = roc_curve(y_test, pred_prob_logit[:, 1], pos_label='Neutral')
fpr[2], tpr[2], thresh[2] = roc_curve(y_test, pred_prob_logit[:, 2], pos_label='Positive')

# plotting
plt.plot(fpr[0], tpr[0], linestyle='--', color='orange', label='Negative vs Rest')
plt.plot(fpr[1], tpr[1], linestyle='--', color='green', label='Neutral vs Rest')
plt.plot(fpr[2], tpr[2], linestyle='--', color='blue', label='Positive vs Rest')
plt.title('Multiclass ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive rate')
plt.legend(loc='best')
```

```
Out[4]: <matplotlib.legend.Legend at 0x7f9f9b7d7430>
```



Penalty=none

```
In [5]: clf_logit_2 = LogisticRegression(penalty='none', dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1,
random_state=None, solver='lbfgs', max_iter=1000, multi_class='auto', verbose=0)
clf_logit_2.fit(X_train_smt, y_train_smt)
pred_logit_2 = clf_logit_2.predict(X_test)
pred_prob_logit_2 = clf_logit_2.predict_proba(X_test)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_ = 1000
```

```
In [6]: print('Predicted classes:', clf_logit_2.classes_)
print('Average accuracy:', np.mean(pred_logit_2 == y_test)*100)
print('Train accuracy:', (clf_logit_2.score(X_train_smt, y_train_smt))*100)
print('Test accuracy:', (clf_logit_2.score(X_test, y_test))*100)

print('\n CONFUSION MATRIX')
print(confusion_matrix(y_test, pred_logit_2))
print('\n CLASSIFICATION REPORT')
print(classification_report(y_test, pred_logit_2))

Predicted classes: ['Negative' 'Neutral' 'Positive']
Average accuracy: 87.7641006777218
Train accuracy: 99.43327419672777
Test accuracy: 87.7641006777218

CONFUSION MATRIX
[[ 37  16  66]
 [ 20  61 163]
 [ 70 325 4679]]

CLASSIFICATION REPORT
              precision    recall  f1-score   support

Negative      0.28         0.31         0.29        119
Neutral       0.15         0.24         0.19        250
Positive      0.95         0.92         0.94       5074


accuracy      0.46         0.49         0.88       5443
macro avg     0.46         0.49         0.47       5443
weighted avg   0.90         0.88         0.89       5443
```

```
In [7]: fpr = {}
tpr = {}
thresh = {}

fpr[0], tpr[0], thresh[0] = roc_curve(y_test, pred_prob_logit_2[:, 0], pos_label='Negative')
fpr[1], tpr[1], thresh[1] = roc_curve(y_test, pred_prob_logit_2[:, 1], pos_label='Neutral')
fpr[2], tpr[2], thresh[2] = roc_curve(y_test, pred_prob_logit_2[:, 2], pos_label='Positive')

# plotting
plt.plot(fpr[0], tpr[0], linestyle='--', color='orange', label='Negative vs Rest')
plt.plot(fpr[1], tpr[1], linestyle='--', color='green', label='Neutral vs Rest')
plt.plot(fpr[2], tpr[2], linestyle='--', color='blue', label='Positive vs Rest')
plt.title('Multiclass ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive rate')
plt.legend(loc='best')
```

```
Out[7]: <matplotlib.legend.Legend at 0x7f9f9b5f910>
```



Solver = newton-cg

```
In [8]: clf_logit_3 = LogisticRegression(penalty='l2', dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1,
random_state=None, solver='newton-cg', max_iter=1000, multi_class='auto', verbose=0)
clf_logit_3.fit(X_train_smt, y_train_smt)
pred_logit_3 = clf_logit_3.predict(X_test)
pred_prob_logit_3 = clf_logit_3.predict_proba(X_test)
```

```
In [9]: print('Predicted classes:', clf_logit_3.classes_)
print('Average accuracy:', np.mean(pred_logit_3 == y_test)*100)
print('Train accuracy:', (clf_logit_3.score(X_train_smt, y_train_smt))*100)
print('Test accuracy:', (clf_logit_3.score(X_test, y_test))*100)

print('\n CONFUSION MATRIX')
print(confusion_matrix(y_test, pred_logit_3))
print('\n CLASSIFICATION REPORT')
print(classification_report(y_test, pred_logit_3))

Predicted classes: ['Negative' 'Neutral' 'Positive']
Average accuracy: 99.02388388756201
Train accuracy: 93.24360400460808
Test accuracy: 99.02388388756201

CONFUSION MATRIX
[[ 49  30  40]
 [ 31 101 118]
 [ 38 280 4759]]

CLASSIFICATION REPORT
              precision    recall  f1-score   support

Negative      0.42         0.41         0.41        119
Neutral       0.15         0.24         0.19        250
Positive      0.95         0.92         0.94       5074

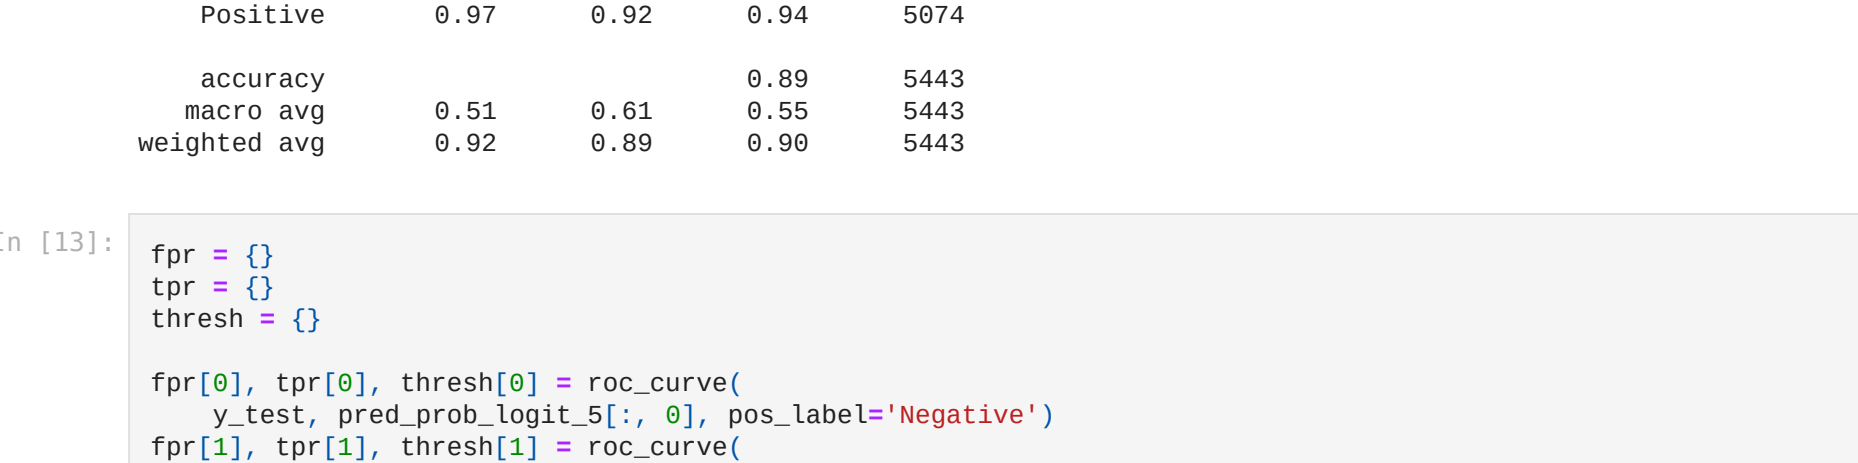
accuracy      0.54         0.58         0.90       5443
macro avg     0.54         0.58         0.56       5443
weighted avg   0.92         0.90         0.91       5443
```

```
In [10]: fpr = {}
tpr = {}
thresh = {}

fpr[0], tpr[0], thresh[0] = roc_curve(y_test, pred_prob_logit_3[:, 0], pos_label='Negative')
fpr[1], tpr[1], thresh[1] = roc_curve(y_test, pred_prob_logit_3[:, 1], pos_label='Neutral')
fpr[2], tpr[2], thresh[2] = roc_curve(y_test, pred_prob_logit_3[:, 2], pos_label='Positive')

# plotting
plt.plot(fpr[0], tpr[0], linestyle='--', color='orange', label='Negative vs Rest')
plt.plot(fpr[1], tpr[1], linestyle='--', color='green', label='Neutral vs Rest')
plt.plot(fpr[2], tpr[2], linestyle='--', color='blue', label='Positive vs Rest')
plt.title('Multiclass ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive rate')
plt.legend(loc='best')
```

```
Out[10]: <matplotlib.legend.Legend at 0x7f9f9b5f9bb0>
```



Penalty=none, solver=newton-cg

```
In [11]: clf_logit_4 = LogisticRegression(penalty='none', dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1,
random_state=None, solver='newton-cg', max_iter=1000, multi_class='auto', verbose=0)
clf_logit_4.fit(X_train_smt, y_train_smt)
pred_logit_4 = clf_logit_4.predict(X_test)
pred_prob_logit_4 = clf_logit_4.predict_proba(X_test)
```

```
In [12]: print('Predicted classes:', clf_logit_4.classes_)
print('Average accuracy:', np.mean(pred_logit_4 == y_test)*100)
print('Train accuracy:', (clf_logit_4.score(X_train_smt, y_train_smt))*100)
print('Test accuracy:', (clf_logit_4.score(X_test, y_test))*100)

print('\n CONFUSION MATRIX')
print(confusion_matrix(y_test, pred_logit_4))
print('\n CLASSIFICATION REPORT')
print(classification_report(y_test, pred_logit_4))

Predicted classes: ['Negative' 'Neutral' 'Positive']
Average accuracy: 99.02388388756201
Train accuracy: 93.24360400460808
Test accuracy: 99.02388388756201

CONFUSION MATRIX
[[ 49  30  40]
 [ 31 101 118]
 [ 38 280 4759]]

CLASSIFICATION REPORT
              precision    recall  f1-score   support

Negative      0.42         0.41         0.41        119
Neutral       0.15         0.24         0.19        250
Positive      0.95         0.92         0.94       5074

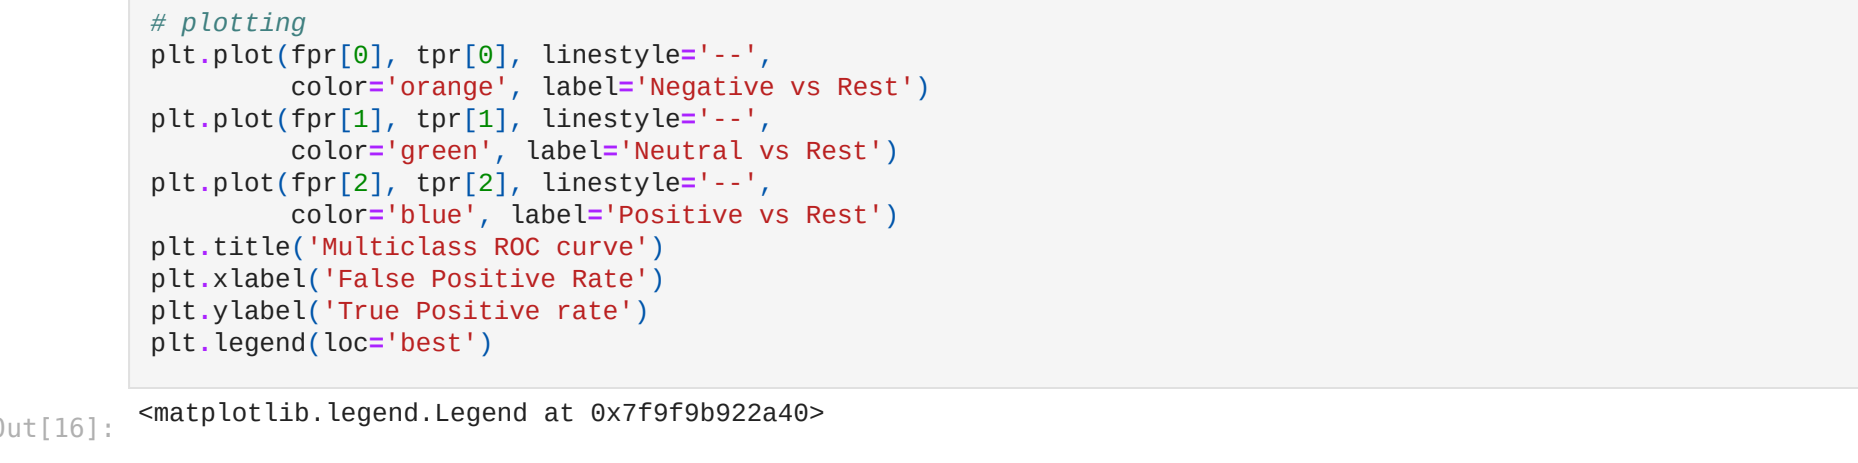
accuracy      0.54         0.58         0.90       5443
macro avg     0.54         0.58         0.56       5443
weighted avg   0.92         0.90         0.91       5443
```

```
In [13]: fpr = {}
tpr = {}
thresh = {}

fpr[0], tpr[0], thresh[0] = roc_curve(y_test, pred_prob_logit_4[:, 0], pos_label='Negative')
fpr[1], tpr[1], thresh[1] = roc_curve(y_test, pred_prob_logit_4[:, 1], pos_label='Neutral')
fpr[2], tpr[2], thresh[2] = roc_curve(y_test, pred_prob_logit_4[:, 2], pos_label='Positive')

# plotting
plt.plot(fpr[0], tpr[0], linestyle='--', color='orange', label='Negative vs Rest')
plt.plot(fpr[1], tpr[1], linestyle='--', color='green', label='Neutral vs Rest')
plt.plot(fpr[2], tpr[2], linestyle='--', color='blue', label='Positive vs Rest')
plt.title('Multiclass ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive rate')
plt.legend(loc='best')
```

```
Out[13]: <matplotlib.legend.Legend at 0x7f9f9b2f5f80>
```



Dual=True, solver=liblinear, multi_class=ovr

```
In [14]: clf_logit_6 = LogisticRegression(penalty='l2', dual=True, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1,
random_state=None, solver='liblinear', max_iter=1000, multi_class='ovr', verbose=0)
clf_logit_6.fit(X_train_smt, y_train_smt)
pred_logit_6 = clf_logit_6.predict(X_test)
pred_prob_logit_6 = clf_logit_6.predict_proba(X_test)
```

```
In [15]: print('Predicted classes:', clf_logit_6.classes_)
print('Average accuracy:', np.mean(pred_logit_6 == y_test)*100)
print('Train accuracy:', (clf_logit_6.score(X_train_smt, y_train_smt))*100)
print('Test accuracy:', (clf_logit_6.score(X_test, y_test))*100)

print('\n CONFUSION MATRIX')
print(confusion_matrix(y_test, pred_logit_6))
print('\n CLASSIFICATION REPORT')
print(classification_report(y_test, pred_logit_6))

Predicted classes: ['Negative' 'Neutral' 'Positive']
Average accuracy: 89.12364040460808
Train accuracy: 93.27813918814507
Test accuracy: 89.12364040460808

CONFUSION MATRIX
[[ 61  26  32]
 [ 36 102 112]
 [ 55 331 4688]]

CLASSIFICATION REPORT
              precision    recall  f1-score   support

Negative      0.49         0.51         0.45        119
Neutral       0.22         0.41         0.29        250
Positive      0.97         0.92         0.95       5074


accuracy      0.53         0.61         0.89       5443
macro avg     0.53         0.61         0.55       5443
weighted avg   0.92         0.89         0.91       5443
```

```
In [16]: fpr = {}
tpr = {}
thresh = {}

fpr[0], tpr[0], thresh[0] = roc_curve(y_test, pred_prob_logit_6[:, 0], pos_label='Negative')
fpr[1], tpr[1], thresh[1] = roc_curve(y_test, pred_prob_logit_6[:, 1], pos_label='Neutral')
fpr[2], tpr[2], thresh[2] = roc_curve(y_test, pred_prob_logit_6[:, 2], pos_label='Positive')

# plotting
plt.plot(fpr[0], tpr[0], linestyle='--', color='orange', label='Negative vs Rest')
plt.plot(fpr[1], tpr[1], linestyle='--', color='green', label='Neutral vs Rest')
plt.plot(fpr[2], tpr[2], linestyle='--', color='blue', label='Positive vs Rest')
plt.title('Multiclass ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive rate')
plt.legend(loc='best')
```

```
Out[16]: <matplotlib.legend.Legend at 0x7f9f9b224a0>
```



Solver=sag

```
In [17]: clf_logit_7 = LogisticRegression(penalty='l2', dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1,
random_state=None, solver='sag', max_iter=1000, multi_class='ovr', verbose=0)
clf_logit_7.fit(X_train_smt, y_train_smt)
pred_logit_7 = clf_logit_7.predict(X_test)
pred_prob_logit_7 = clf_logit_7.predict_proba(X_test)
```

```
In [18]: print('Predicted classes:', clf_logit_7.classes_)
print('Average accuracy:', np.mean(pred_logit_7 == y_test)*100)
print('Train accuracy:', (clf_logit_7.score(X_train_smt, y_train_smt))*100)
print('Test accuracy:', (clf_logit_7.score(X_test, y_test))*100)

print('\n CONFUSION MATRIX')
print(confusion_matrix(y_test, pred_logit_7))
print('\n CLASSIFICATION REPORT')
print(classification_report(y_test, pred_logit_7))

Predicted classes: ['Negative' 'Neutral' 'Positive']
Average accuracy: 90.13411721477127
Train accuracy: 91.008420709094
Test accuracy: 90.13411721477127

CONFUSION MATRIX
[[ 50  30  39]
 [ 31 102 117]
 [ 48 280 4751]]

CLASSIFICATION REPORT
              precision    recall  f1-score   support

Negative      0.41         0.42         0.42        119
Neutral       0.25         0.41         0.31        250
Positive      0.97         0.94         0.95       5074


accuracy      0.54         0.59         0.90       5443
macro avg     0.54         0.59         0.56       5443
weighted avg   0.92         0.90         0.91       5443
```

```
In [19]: fpr = {}
tpr = {}
thresh = {}

fpr[0], tpr[0], thresh[0] = roc_curve(y_test, pred_prob_logit_7[:, 0], pos_label='Negative')
fpr[1], tpr[1], thresh[1] = roc_curve(y_test, pred_prob_logit_7[:, 1], pos_label='Neutral')
fpr[2], tpr[2], thresh[2] = roc_curve(y_test, pred_prob_logit_7[:, 2], pos_label='Positive')

# plotting
plt.plot(fpr[0], tpr[0], linestyle='--', color='orange', label='Negative vs Rest')
plt.plot(fpr[1], tpr[1], linestyle='--', color='green', label='Neutral vs Rest')
plt.plot(fpr[2], tpr[2], linestyle='--', color='blue', label='Positive vs Rest')
plt.title('Multiclass ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive rate')
plt.legend(loc='best')
```

```
Out[19]: <matplotlib.legend.Legend at 0x7f9f9f75d28>
```



Penalty=none, solver=sag

```
In [20]: clf_logit_8 = LogisticRegression(penalty='l2', dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1,
random_state=None, solver='sag', max_iter=1000, multi_class='ovr', verbose=0)
clf_logit_8.fit(X_train_smt, y_train_smt)
pred_logit_8 = clf_logit_8.predict(X_test)
pred_prob_logit_8 = clf_logit_8.predict_proba(X_test)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn("The max_iter was reached which means the coef_ did not converge")
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn("The max_iter was reached which means the coef_ did not converge")
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn("The max_iter was reached which means the coef_ did not converge")
```

```
In [21]: print('Predicted classes:', clf_logit_8.classes_)
print('Average accuracy:', np.mean(pred_logit_8 == y_test)*100)
print('Train accuracy:', (clf_logit_8.score(X_train_smt, y_train_smt))*100)
print('Test accuracy:', (clf_logit_8.score(X_test, y_test))*100)

print('\n CONFUSION MATRIX')
print(confusion_matrix(y_test, pred_logit_8))
print('\n CLASSIFICATION REPORT')
print(classification_report(y_test, pred_logit_8))

Predicted classes: ['Negative' 'Neutral' 'Positive']
Average accuracy: 88.62759507624472
Train accuracy: 90.6180250808404
Test accuracy: 88.62759507624472

CONFUSION MATRIX
[[ 30  22  67]
 [ 26  63 161]
 [ 63 280 4731]]

CLASSIFICATION REPORT
              precision    recall  f1-score   support

Negative      0.25         0.25         0.25        119
Neutral       0.17         0.25         0.20        250
Positive      0.95         0.93         0.94       5074

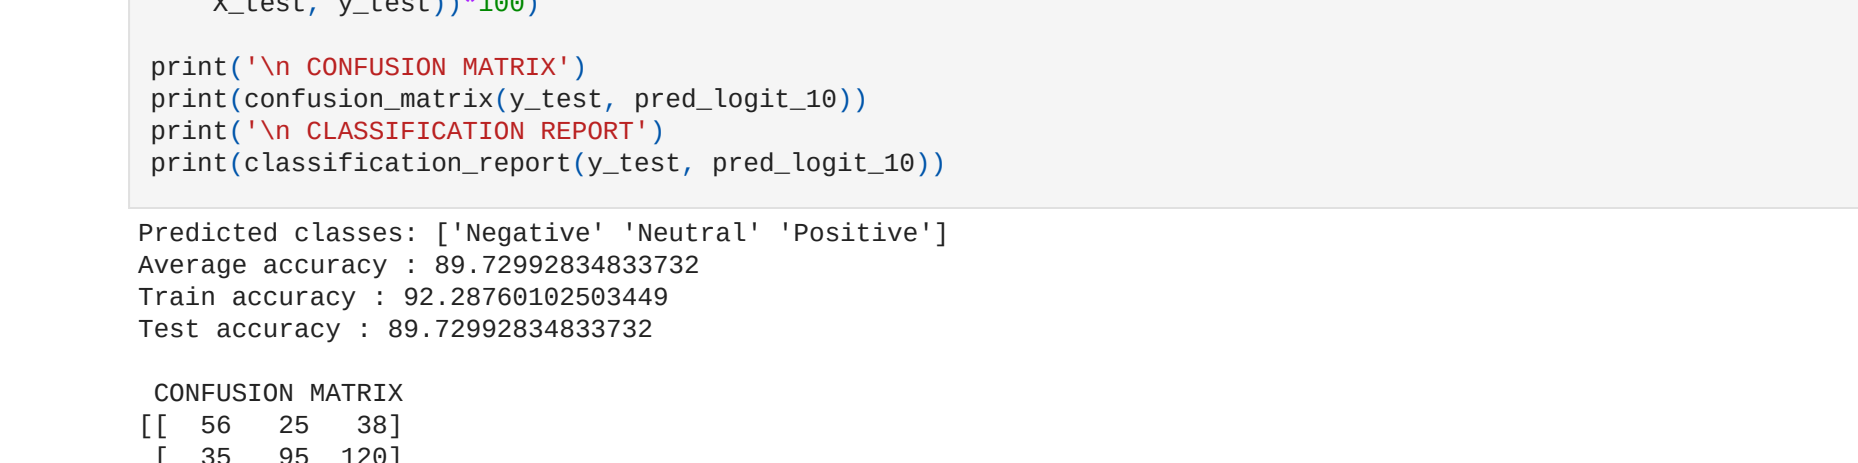
accuracy      0.46         0.48         0.87       5443
macro avg     0.46         0.48         0.47       5443
weighted avg   0.90         0.89         0.89       5443
```

```
In [22]: fpr = {}
tpr = {}
thresh = {}

fpr[0], tpr[0], thresh[0] = roc_curve(y_test, pred_prob_logit_8[:, 0], pos_label='Negative')
fpr[1], tpr[1], thresh[1] = roc_curve(y_test, pred_prob_logit_8[:, 1], pos_label='Neutral')
fpr[2], tpr[2], thresh[2] = roc_curve(y_test, pred_prob_logit_8[:, 2], pos_label='Positive')

# plotting
plt.plot(fpr[0], tpr[0], linestyle='--', color='orange', label='Negative vs Rest')
plt.plot(fpr[1], tpr[1], linestyle='--', color='green', label='Neutral vs Rest')
plt.plot(fpr[2], tpr[2], linestyle='--', color='blue', label='Positive vs Rest')
plt.title('Multiclass ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive rate')
plt.legend(loc='best')
```

```
Out[22]: <matplotlib.legend.Legend at 0x7f9f9f7136a0>
```



Penalty=l1, solver=saga

```
In [26]: clf_logit_10 = LogisticRegression(penalty='l1', dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1,
random_state=None, solver='saga', max_iter=1000, multi_class='ovr', verbose=0)
clf_logit_10.fit(X_train_smt, y_train_smt)
pred_logit_10 = clf_logit_10.predict(X_test)
pred_prob_logit_10 = clf_logit_10.predict_proba(X_test)
```

```
In [27]: print('Predicted classes:', clf_logit_10.classes_)
print('Average accuracy:', np.mean(pred_logit_10 == y_test)*100)
print('Train accuracy:', (clf_logit_10.score(X_train_smt, y_train_smt))*100)
print('Test accuracy:', (clf_logit_10.score(X_test, y_test))*100)

print('\n CONFUSION MATRIX')
print(confusion_matrix(y_test, pred_logit_10))
print('\n CLASSIFICATION REPORT')
print(classification_report(y_test, pred_logit_10))

Predicted classes: ['Negative' 'Neutral' 'Positive']
Average accuracy: 89.72992834833732
Train accuracy: 92.25765027059491
Test accuracy: 89.72992834833732

CONFUSION MATRIX
[[ 56  25  38]
 [ 35  85 122]
 [ 44 297 4733]]

CLASSIFICATION REPORT
              precision    recall  f1-score   support

Negative      0.45         0.47         0.44        119
Neutral       0.23         0.38         0.29        250
Positive      0.97         0.93         0.95       5074

accuracy      0.54         0.59         0.90       5443
macro avg     0.54         0.59         0.56       5443
weighted avg   0.92         0.90         0.91       5443
```

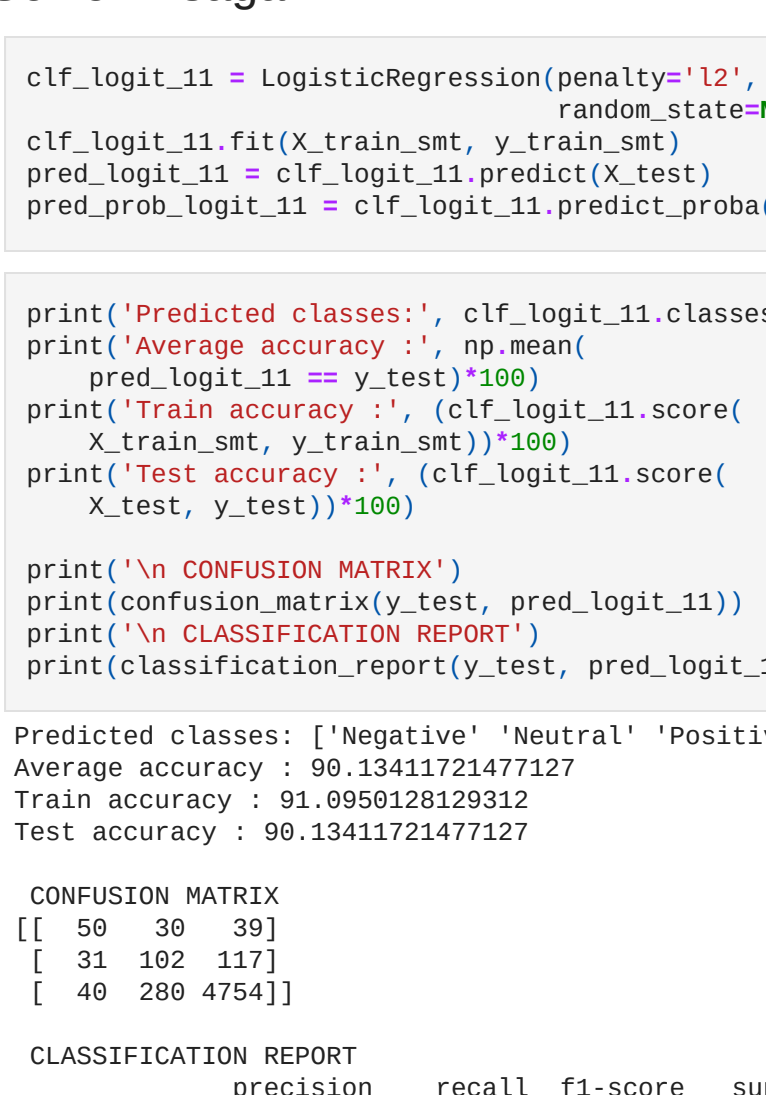


```
In [29]: fpr = {}
tpr = {}
thresh = {}

fpr[0], tpr[0], thresh[0] = roc_curve(
    y_test, pred_prob_logit_10[:, 0], pos_label='Negative')
fpr[1], tpr[1], thresh[1] = roc_curve(
    y_test, pred_prob_logit_10[:, 1], pos_label='Neutral')
fpr[2], tpr[2], thresh[2] = roc_curve(
    y_test, pred_prob_logit_10[:, 2], pos_label='Positive')

# plotting
plt.plot(fpr[0], tpr[0], linestyle='--',
         color='orange', label='Negative vs Rest')
plt.plot(fpr[1], tpr[1], linestyle='--',
         color='green', label='Neutral vs Rest')
plt.plot(fpr[2], tpr[2], linestyle='--',
         color='blue', label='Positive vs Rest')
plt.title('Multiclass ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive rate')
plt.legend(loc='best')
```

Out[28]: <matplotlib.legend.Legend at 0x7f9f9162f550>



Solver = saga

```
In [29]: clf_logit_11 = LogisticRegression(penalty='l2', dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_
pred_prob_logit_11 = clf_logit_11.predict(X_test)
```

```
In [30]: print('Predicted classes:', clf_logit_11.classes_)
print('Average accuracy:', np.mean(
    pred_logit_11 == y_test)*100)
print('Train accuracy:', (clf_logit_11.score(
    X_train_smt, y_train_smt))*100)
print('Test accuracy:', (clf_logit_11.score(
    X_test, y_test))*100)

print('\n CONFUSION MATRIX')
print(confusion_matrix(y_test, pred_logit_11))
print('\n CLASSIFICATION REPORT')
print(classification_report(y_test, pred_logit_11))

Predicted classes: ['Negative' 'Neutral' 'Positive']
Average accuracy : 90.13411721477127
Train accuracy : 91.9909129129512
Test accuracy : 90.13411721477127

CONFUSION MATRIX
[[ 50  30  39]
 [ 31 102 117]
 [ 49 298 475]]

CLASSIFICATION REPORT
              precision    recall  f1-score   support

Negative      0.41         0.42         0.42         119
Neutral       0.25         0.41         0.31         250
Positive      0.97         0.94         0.95         5074

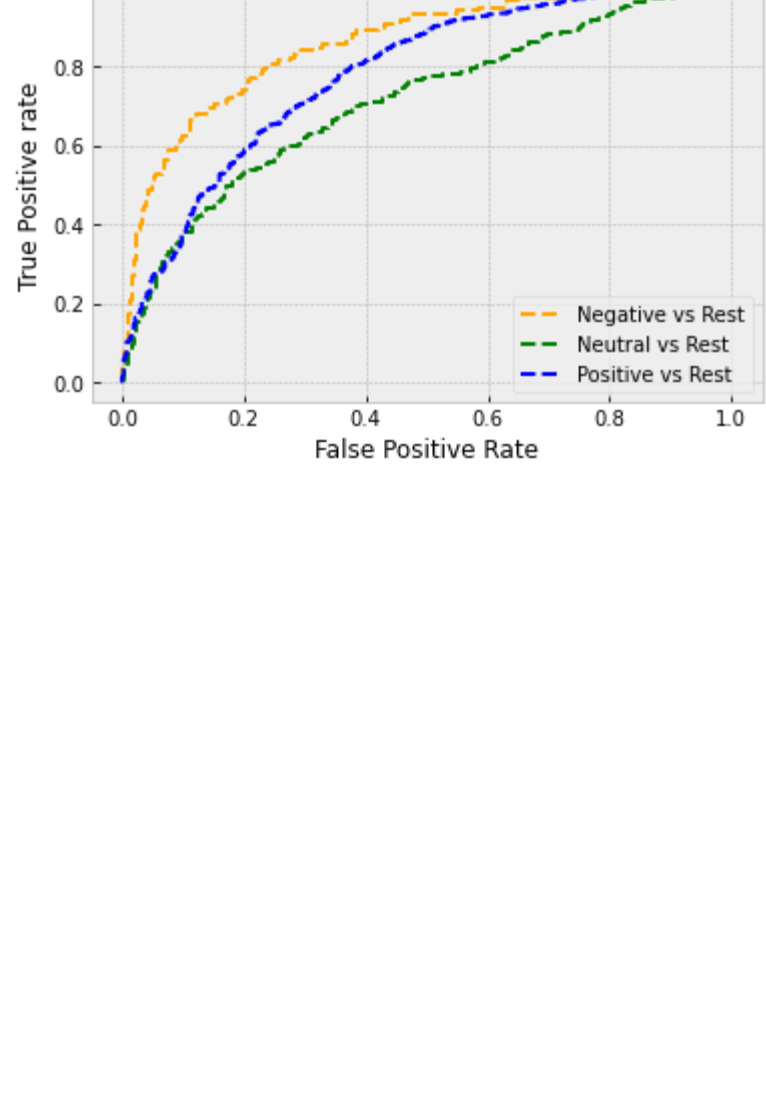
accuracy              0.90         0.90         0.90         5443
macro avg          0.54         0.59         0.56         5443
weighted avg       0.92         0.90         0.91         5443
```

```
In [31]: fpr = {}
tpr = {}
thresh = {}

fpr[0], tpr[0], thresh[0] = roc_curve(
    y_test, pred_prob_logit_11[:, 0], pos_label='Negative')
fpr[1], tpr[1], thresh[1] = roc_curve(
    y_test, pred_prob_logit_11[:, 1], pos_label='Neutral')
fpr[2], tpr[2], thresh[2] = roc_curve(
    y_test, pred_prob_logit_11[:, 2], pos_label='Positive')

# plotting
plt.plot(fpr[0], tpr[0], linestyle='--',
         color='orange', label='Negative vs Rest')
plt.plot(fpr[1], tpr[1], linestyle='--',
         color='green', label='Neutral vs Rest')
plt.plot(fpr[2], tpr[2], linestyle='--',
         color='blue', label='Positive vs Rest')
plt.title('Multiclass ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive rate')
plt.legend(loc='best')
```

Out[31]: <matplotlib.legend.Legend at 0x7f9f916c32e0>



Penalty = none

```
In [32]: clf_logit_12 = LogisticRegression(penalty='none', dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_
pred_logit_12 = clf_logit_12.predict(X_test)
pred_prob_logit_12 = clf_logit_12.predict_proba(X_test)
```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge

```
In [33]: print('Predicted classes:', clf_logit_12.classes_)
print('Average accuracy:', np.mean(
    pred_logit_12 == y_test)*100)
print('Train accuracy:', (clf_logit_12.score(
    X_train_smt, y_train_smt))*100)
print('Test accuracy:', (clf_logit_12.score(
    X_test, y_test))*100)

print('\n CONFUSION MATRIX')
print(confusion_matrix(y_test, pred_logit_12))
print('\n CLASSIFICATION REPORT')
print(classification_report(y_test, pred_logit_12))

Predicted classes: ['Negative' 'Neutral' 'Positive']
Average accuracy : 88.8296895994617
Train accuracy : 88.9382646691635
Test accuracy : 88.8296895994617

CONFUSION MATRIX
[[ 28  23  68]
 [ 23  64 163]
 [ 66 265 474]]

CLASSIFICATION REPORT
              precision    recall  f1-score   support

Negative      0.24         0.24         0.24         119
Neutral       0.18         0.26         0.21         250
Positive      0.95         0.93         0.94         5074

accuracy              0.89         0.89         0.89         5443
macro avg          0.46         0.48         0.46         5443
weighted avg       0.90         0.89         0.90         5443
```

```
In [34]: fpr = {}
tpr = {}
thresh = {}

fpr[0], tpr[0], thresh[0] = roc_curve(
    y_test, pred_prob_logit_12[:, 0], pos_label='Negative')
fpr[1], tpr[1], thresh[1] = roc_curve(
    y_test, pred_prob_logit_12[:, 1], pos_label='Neutral')
fpr[2], tpr[2], thresh[2] = roc_curve(
    y_test, pred_prob_logit_12[:, 2], pos_label='Positive')

# plotting
plt.plot(fpr[0], tpr[0], linestyle='--',
         color='orange', label='Negative vs Rest')
plt.plot(fpr[1], tpr[1], linestyle='--',
         color='green', label='Neutral vs Rest')
plt.plot(fpr[2], tpr[2], linestyle='--',
         color='blue', label='Positive vs Rest')
plt.title('Multiclass ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive rate')
plt.legend(loc='best')
```

Out[34]: <matplotlib.legend.Legend at 0x7f9f9153a629>

